

Delivered but Not Heard: AI Proxy Speech in Power-Imbalanced Group Decision-Making*

Soohwan Lee
Department of Design
UNIST
Republic of Korea
soohwanlee@unist.ac.kr

Kyungho Lee
Department of Design
UNIST
Republic of Korea
kyungho@unist.ac.kr

Abstract

AI-mediated communication (AIMC) can relay a person's concern under an AI identity. We use the term *AI proxy speech* for this arrangement. We revisit our study of 96 participants in groups with unequal power in a controlled setting. The AI relayed junior members' private input under its own name alongside counterarguments it generated. Some senior members discounted AI messages, and juniors reported lower psychological safety in this condition than in the baseline. The design also limited timely clarification and follow-up. We argue that evaluating AI proxy speech requires attention to message delivery, group response, and opportunities for protected follow-up. Beyond this controlled setting, cultural and organizational conditions may create needs for protected communication that warrant further study. We propose examining when such channels could help members raise concerns while limiting personal exposure, how groups respond, and how those who hold authority shape adoption and use.

CCS Concepts

• **Human-centered computing** → **Computer supported cooperative work**; *Collaborative interaction*; *Natural language interfaces*; *HCI theory, concepts and models*.

Keywords

AI-mediated communication, AI proxy speech, minority voice, group decision-making, power imbalance, psychological safety

1 Introduction

Groups can make better decisions when members bring different information and question each other's reasoning. The same setting can also suppress the views it needs. Members who disagree may voice agreement anyway, a response social psychology calls compliance: a public position that does not match a private one [1, 24, 33]. A member with a minority view faces pressure from a majority whose agreement is already visible. A member with less power may also depend on others for evaluation, rewards, and future work [7, 8]. These pressures can lead members to withhold concerns that could inform the group's decision [22]. Speaking identifies the person who disagrees and may expose them to consequences beyond the discussion.

AI-mediated communication (AIMC) changes the link between a person and a message. It covers communication in which AI

changes, adds to, or generates a message for a communicator [13]. Many systems keep the person as the visible sender, but AI can also appear as the public source. A junior employee who fears retaliation for reporting a safety risk might want this distance from the message. In unsuccessful conversations, participants assigned less responsibility to their partner when using a messenger with AI reply suggestions [16]. This raises a related question for group communication: how do members respond when AI presents someone's concern under its own identity?

We use *AI proxy speech* as a descriptive term for an AIMC setting in which a person supplies the concern while AI appears as the public speaker. We argue that evaluating AI proxy speech requires attention to message delivery, group response, and opportunities for protected follow-up. We develop this position by revisiting our completed study of minority support in groups with unequal power [27], where the channel did not provide the support we intended.

This paper offers the workshop three things. It contributes a case in which an AI-mediated channel delivered a member's concern under an AI identity, yet some members gave AI messages less weight. It raises a design question about how to limit exposure of the sender while supporting clarification and follow-up. It also raises a governance question, since people whose authority creates a need for protected communication may also decide whether such a channel exists. We hope to discuss how CSCW should evaluate and govern proxy speech for members with less power.

2 Related Work

2.1 Power, Silence, and the Limits of Anonymity

Social influence research explains why people who hold a different view often keep it private. Majority opinion can produce compliance, a public agreement that does not reflect private belief [1, 33]. Kelman separates compliance from real attitude change and links it to avoiding punishment or seeking reward [24]. Formal hierarchy can reinforce this pressure when senior members hold legitimate and reward power over junior members [7, 8]. When the minority stays silent, the group loses the views that could improve its decision, which is a known path to groupthink [22]. A requirement for unanimous agreement can also reinforce commitment to an earlier decision [23]. Organizations can also develop shared beliefs that raising problems is unsafe [32]. Psychological safety can counter this, but it depends on team and leadership practice [6].

Two established remedies show why the problem is hard. The first uses a devil's advocate to challenge the emerging view, a role also explored with AI-generated counterarguments [26]. This can encourage discussion of different perspectives [30], but the assigned advocate may not hold the view they present. Prior work finds that

*Position paper accepted to the workshop Broader Impacts of GenAI in Communication: Building Agendas for Research, Design, and Policymaking at the 2026 ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW '26), Salt Lake City, UT, USA, October 10, 2026.

genuine dissent can prompt broader consideration of alternatives than assigned dissent [35, 42]. The assigned role also offers no direct access to concerns that others keep private. The second remedy removes the name from the message, which can lower the risk of exposure but changes how members identify with the group and how they express criticism [2, 39]. One AI-based approach relays a member's private input under the AI's identity [29]. Anonymity has two limits that matter here. In a small group, a unique opinion may still identify its source; research on whistleblowing also shows that content and writing style can enable re-identification [45]. Removing identifying details also leaves open how a group will respond to the message. Protection and uptake therefore require separate attention.

2.2 Human Input, Public Attribution, and Responsibility in AIMC

Many AIMC systems keep the person as the visible speaker while suggesting replies, changing tone, or helping someone present themselves [9, 11, 15, 17, 40]. Message suggestions can also support cooperation between people with different group identities [4]. Support can sit on the receiving side, where AI reframes a difficult message for the person who reads it [49]. Group agents can prompt participation [5], help learners form social connections [47], mediate between views [43, 46], or generate counterarguments [3]. People and systems may share topic selection, wording, and approval across these arrangements. Our case combines the AI's own counterarguments with messages based on a member's private input. We use *proxy speech* for the latter: the AI paraphrases a concern and posts it under its own identity. Receivers see the same public source for both types of message. Work on agents that represent people through generated speech also raises questions about user control and trust [19].

AI involvement can affect both a message and how people assign responsibility for it. Suggested text can change topic choice and expressed views [20, 36]. More support also does not always give users more agency. One study of personalized writing support found more reading highlights alongside lower reported autonomy, ownership, and self-credit [38]. AIMC studies also examine changes in language use and interpersonal perceptions [17, 31]. Beliefs about AI authorship can affect judgments of a communicator [21], and AI-generated replies can affect responsibility judgments after a conversation fails [16]. These findings show why evaluation should distinguish message production, perceived authorship, and responsibility.

Recent disclosure studies also report mixed results. One experiment found lower perceived author trustworthiness and authenticity, and lower willingness to use the content, when AI involvement was disclosed [41]. Another study found lower judgments of writers after disclosure, especially for social and emotional writing [34], while two trust game experiments found no significant reduction in behavioral or stated trust after disclosure [37]. A protected channel adds a further problem. Marking a specific message as based on human input may reveal that someone in a small group disagrees. Labeling it only as AI generated may hide that a member has something at stake. Disclosure choices therefore require attention to both source visibility and the risk of identifying the sender.

3 Case Study: AI-Mediated Communication for Minority Voice

3.1 Study Design and Intervention

Our prior study involved 96 Korean participants in 24 groups of four [27]. Each group included three senior members with high power and one junior member with low power. Role labels appeared beside messages. Seniors jointly decided whether the junior received a bonus, giving them evaluation and reward authority. Groups discussed two workplace decisions through text chat: selecting a team leader for promotion and choosing a contractor. Task materials gave seniors and juniors reasons to prefer different options, placing the junior in a minority position.

We compared a baseline without AI support with two interventions. In the AI-generated Counterargument condition (AIGC, Figure 1, Case 1), the AI posted arguments against the emerging majority view. In the AI-mediated Message condition (AIMM, Case 2), it also received private input from the junior, paraphrased it, and posted it under the AI's name. AIMM thus combined AI-generated counterarguments with messages based on human input under the same identity. This mixture was intended to protect the junior's anonymity by concealing the source of each post. Each participant experienced the baseline and one AI condition. Only the junior knew about the private channel. Seniors therefore saw similar AI posts in both conditions, but juniors in AIMM knew that some posts represented their own concerns.

3.2 Summary of Empirical Findings

We measured psychological safety, perceived influence, marginalization, satisfaction, and communication volume. We also interviewed juniors individually and seniors as a group. Across conditions, about 80% of final decisions aligned with seniors' preferences. AIGC reduced juniors' reported marginalization and improved several satisfaction measures without a significant reduction in psychological safety. In AIMM, juniors reported lower psychological safety, perceived influence, and process and outcome satisfaction, and greater marginalization than in the baseline. The regression model showed that juniors wrote more characters in AIMM than in the baseline, but this difference did not remain significant after Bonferroni correction. Message counts did not differ significantly. These results provide limited evidence of increased communication volume.

Interviews help explain these experiences. Some seniors gave AI messages less weight because they did not come from a person. Others questioned the AI's credibility because its position changed during the discussion. Juniors sometimes felt ignored when the group dismissed messages that carried their input. These accounts suggest a problem with how receivers interpreted the source, but they do not establish why each relayed message received a particular response.

The system also limited timely clarification and follow-up. It posted at eight-turn intervals and sought to avoid repeating content. Juniors could submit private input, but the design did not support direct exchanges with the AI in the group discussion. Publicly claiming a relayed message would reveal its source. We interpret these constraints as a reason to examine follow-up alongside initial delivery. The interviews also raised questions about authenticity,




Figure 1: The two AI conditions. AIGC posts counterarguments generated by the system. AIMM also paraphrases junior members' private input and presents it under the AI's name. The junior's input and the AI's own arguments appear under the same identity. Junior members reported more favorable experiences in AIGC than in AIMM.

authorship, and responsibility. We use these concepts to interpret participants' accounts; the study did not directly measure them.

4 Implications

4.1 Evaluating Protected Channels by Group Response

Group Experience (GX) offers a lens for examining experiences that emerge through group interaction [28]. For AI proxy speech, this means examining how the group responds as well as how the sender uses the system. A protected channel changes what receivers know, how a sender can respond, and who makes the decision. We consider four questions for design and evaluation. The first concerns what members know about the channel: that it can carry a person's concern, how AI changes the input, and who can access the records. The second concerns the sender's control. A person should be able to review what the AI changed and choose whether to edit, approve, delay, or withdraw the message. The third concerns the group's response. Before reaching a decision, members could acknowledge relevant concerns and record how they considered them, including reasons for rejecting a suggestion. The fourth concerns clarification. A sender may need to correct the AI's wording, add context, or respond to an objection without revealing their identity. Other members may need evidence or an explanation when a claim affects them. The channel should support these exchanges while allowing the sender to decide whether to disclose more information.

The need for a response continues after delivery. A group may discount a concern because AI posted it and proceed with its preferred decision, leaving assumptions unexamined. Research on groupthink connects such patterns to poor decisions [22]. Preserving anonymity alone may leave this problem unresolved. In our case, some seniors discounted AI posts while juniors understood those posts as attempts to express their views. Receivers may overlook the need to respond when no person publicly claims the concern; whether repeated use reinforces this pattern requires further study.

Follow-up also involves the sender. Using a protected channel to insist on a claim while avoiding questions or responsibility for its consequences could be irresponsible. We propose that senders take responsibility for the content they approve and have a way to explain or correct it without public exposure. Groups would still need to assess its accuracy. The same concealment can make it difficult to claim credit for a useful contribution. If a group adopts a suggestion as the AI's idea, the person who supplied it may remain unrecognized. Researchers should examine whether senders can choose to claim their contributions and what risks that choice carries.

Supporting these exchanges would require responsibilities for the sender, system operator, and group. Operators could limit access to identifying records, and groups could review how they handled concerns. An independent facilitator might verify identity or review unanswered concerns, although that access could deter use or create risk in some settings. A protected channel should support efforts to improve direct communication, psychological safety, and responsible leadership. In addition to comparing message volume, researchers could examine whether the group replied, whether the concern entered the decision record, whether the sender could follow up, and whether repeated use exposed the sender. These measures would help assess what happens after a message reaches the group.

4.2 The Need for Protected Communication

Some groups punish dissent, concentrate decision authority, and offer few routes to raise concerns outside the group. In such settings, a protected channel may offer one of the few alternatives to silence. Our results do not establish that AI proxy speech provides effective support under these conditions. They leave open whether different designs could help people whose concerns would otherwise remain private. This possibility warrants research with safeguards appropriate to the risks members face. We built the channel for a setting where we expected it to help. Seniors controlled the junior's bonus, and task materials placed the junior in a minority position.

The controlled study ended after a single meeting. Participants did not risk an ongoing job or working relationship. In an ongoing group, a member may depend on the same people for evaluation and continued membership, lack an independent route to appeal, or expect lasting consequences for dissent. These constraints could create a need for communication that limits direct attribution. The less favorable experiences in our study remain a reason for caution when considering this possibility.

Three conditions could guide further investigation: members face a credible risk of retaliation, exclusion, or unfair evaluation; the concern affects a consequential decision; and the sender chooses the channel after reviewing its protections and limits. Research would need to test whether these conditions make the channel useful. Studies would need to assess whether the system preserves the intended meaning, how others might infer the source, and who can access private input and records. The sender should be able to review each message and decide whether to proceed. Managing the channel could also add work. Some juniors in our study described difficulty coordinating private input with the public discussion. Further studies should examine that burden alongside the group's response and the sender's experience. Comparisons should include available alternatives, such as independent AI counterarguments, drafting support, or communication with a responsible office outside the group. Independent counterarguments may be sufficient when the group needs another perspective; relaying private input may matter when a member holds information the AI cannot supply. The relative value of these options would depend on local relationships, access to support, and the concern itself.

4.3 Cultural Context and Organizational Adoption

Our participants were Korean, and cultural context may have shaped their experience. A single sample offers little basis for identifying that influence. National cultural dimensions describe broad patterns across countries [14]. Prior work links employee voice and silence to both cultural context and local conditions [18, 25]. Research also distinguishes speaking to suggest improvements from speaking to warn about problems [48]. These findings suggest that the reasons for using a protected channel may vary with the concern, the group, and the cultural setting. Making disagreement visible and resolving it through open discussion may not suit every cultural setting. Individualistic contexts often favor addressing conflict directly, whereas collectivistic contexts may favor indirect communication and norms that preserve interpersonal harmony [7, 12]. Groups also develop shared practices for handling differences, including rules that maintain relationships and limit overt conflict [10]. Conflict resolution may therefore involve little public confrontation. These are broad tendencies, and preferences can vary within cultures and organizations. Researchers should consider when members prefer direct discussion, private mediation, or indirect communication, and whether each approach allows a concern to receive a fair response. A protected channel may fit some of these preferences, but studies would need to examine its effects on both the sender and the group.

Organizational adoption raises a further question. People whose authority creates a need for protected communication may also

decide whether a channel exists. They may accept AI advice while resisting a route for unnamed concerns. They could restrict who uses the channel, demand access to identifying records, or treat posting a message as sufficient participation without considering it. Closed organizational conditions may increase the need for protection while preventing adoption or weakening safeguards. Attending to these practical barriers could be as useful as studying effective interventions in controlled settings. Such research could examine who can block adoption, whose concerns enter the decision, and whether members with less power can influence the terms of use.

Researchers could follow selection, configuration, use, and review of a channel. Possible safeguards include consulting members with less power, limiting access to identifying records, and reviewing how managers respond to concerns. Work on organizational oversight offers a basis for considering these arrangements [44]. Their value would depend on who conducts the review and whether members can question its decisions. Comparative field studies should also examine decision authority, penalties for dissent, access to appeal, and dependence on others for evaluation. These conditions may explain differences within cultural settings as well as across them and help identify when a channel offers meaningful protection.

5 Conclusion

In our study, junior members reported less favorable experiences when the AI relayed their concerns under its own name. Some seniors discounted AI messages, and the design limited timely follow-up. We therefore propose evaluating message delivery, group response, and opportunities to clarify a concern without revealing the sender. Complex organizational and cultural settings may create needs that this controlled study did not capture. Establishing whether AI proxy speech can address them requires further research into both the channel and the people who decide whether and how it is used.

Acknowledgments

The authors gratefully acknowledge Dr. Angel Hsing-Chi Hwang and Dr. Oh-Sang Kwon for their assistance with experimental design and data analysis. This research was partially supported by a grant from the Korea Institute for Advancement of Technology (KIAT) funded by the Government of Korea (MOTIE) (P0025495, Establishment of Infrastructure for Integrated Utilization of Design Industry Data) and the Industrial Technology Innovation Program (RS-2025-25457764, Smart CMF Design Recommendation and Real-Time Collaborative 3D Simulation System) funded by the Ministry of Trade, Industry & Energy (MOTIE, Korea).

References

- [1] Solomon E. Asch. 1955. Opinions and Social Pressure. <https://www.scientificamerican.com/article/opinions-and-social-pressure/>.
- [2] Mikyeong Bae. 2016. The Effects of Anonymity on Computer-Mediated Communication: The Case of Independent versus Interdependent Self-Construal Influence. *Computers in Human Behavior* 55 (Feb. 2016), 300–309. doi:10.1016/j.chb.2015.09.026
- [3] Chun-Wei Chiang, Zhuoran Lu, Zhuoyan Li, and Ming Yin. 2024. Enhancing AI-Assisted Group Decision Making through LLM-Powered Devil's Advocate. In *Proceedings of the 29th International Conference on Intelligent User Interfaces (IUI '24)*. Association for Computing Machinery, New York, NY, USA, 103–119. doi:10.1145/3640543.3645199

- [4] Elijah L. Claggett, Robert E. Kraut, and Hirokazu Shirado. 2025. Relational AI: Facilitating Intergroup Cooperation with Socially Aware Conversational Support. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, 1–22. doi:10.1145/3706598.3713757
- [5] Hyo Jin Do, Ha-Kyung Kong, Jaewook Lee, and Brian P. Bailey. 2022. How Should the Agent Communicate to the Group? Communication Strategies of a Conversational Agent in Group Chat Discussions. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (Nov. 2022), 387:1–387:23. doi:10.1145/3555112
- [6] Amy Edmondson. 1999. Psychological Safety and Learning Behavior in Work Teams. *Administrative Science Quarterly* 44, 2 (June 1999), 350–383. doi:10.2307/2666999
- [7] Donelson R. Forsyth. 2018. *Group Dynamics*. Cengage Learning.
- [8] John R. P. French Jr. and Bertram Raven. 1959. The Bases of Social Power. In *Studies in Social Power*. Univer. Michigan, Oxford, England, 150–167.
- [9] Yue Fu, Sami Foell, Xuhai Xu, and Alexis Hiniker. 2024. From Text to Self: Users' Perception of AIMC Tools on Interpersonal Communication and Self. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/3613904.3641955
- [10] Michele J. Gelfand, Lisa M. Leslie, and Kirsten M. Keller. 2008. On the Etiology of Conflict Cultures. *Research in Organizational Behavior* 28 (2008), 137–166. doi:10.1016/j.riob.2008.06.001
- [11] Kaely Hall, Max Ohsawa, Vedant Das Swain, and Jennifer G Kim. 2026. LLMe: Supporting Identity Work through Human-AI Alignment. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI '26)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/3772318.3790890
- [12] Mitchell R. Hammer. 2005. The Intercultural Conflict Style Inventory: A Conceptual Framework and Measure of Intercultural Conflict Resolution Approaches. *International Journal of Intercultural Relations* 29, 6 (2005), 675–695. doi:10.1016/j.ijintrel.2005.08.010
- [13] Jeffrey T Hancock, Mor Naaman, and Karen Levy. 2020. AI-Mediated Communication: Definition, Research Agenda, and Ethical Considerations. *Journal of Computer-Mediated Communication* 25, 1 (March 2020), 89–100. doi:10.1093/jcmc/zmz022
- [14] Geert Hofstede. 2011. Dimensionalizing Cultures: The Hofstede Model in Context. *Online Readings in Psychology and Culture* 2, 1 (Dec. 2011). doi:10.9707/2307-0919.1014
- [15] Jess Hohenstein and Malte Jung. 2018. AI-Supported Messaging: An Investigation of Human-Human Text Conversation with AI Support. In *Extended Abstracts of the 2018 CHI Conference on Human Factors in Computing Systems (CHI EA '18)*. Association for Computing Machinery, New York, NY, USA, 1–6. doi:10.1145/3170427.3188487
- [16] Jess Hohenstein and Malte Jung. 2020. AI as a Moral Crumple Zone: The Effects of AI-mediated Communication on Attribution and Trust. *Computers in Human Behavior* 106 (May 2020), 106190. doi:10.1016/j.chb.2019.106190
- [17] Jess Hohenstein, Rene F. Kizilcec, Dominic DiFranzo, Zhila Aghajari, Hannah Mieczkowski, Karen Levy, Mor Naaman, Jeffrey Hancock, and Malte F. Jung. 2023. Artificial Intelligence in Communication Impacts Language and Social Relationships. *Scientific Reports* 13, 1 (April 2023), 5487. doi:10.1038/s41598-023-30938-9
- [18] Xu Huang, Evert Van de Vliert, and Gerben Van der Vegt. 2005. Breaking the Silence Culture: Stimulation of Participation and Employee Opinion Withholding Cross-nationally. *Management and Organization Review* 1, 3 (Nov. 2005), 459–482. doi:10.1111/j.1740-8784.2005.00023.x
- [19] Angel Hsing-Chi Hwang, John Oliver Siy, Renee Shelby, and Alison Lentz. 2024. In Whose Voice?: Examining AI Agent Representation of People in Social Interaction through Generative Speech. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference (DIS '24)*. Association for Computing Machinery, New York, NY, USA, 224–245. doi:10.1145/3643834.3661555
- [20] Maurice Jakesch, Advait Bhat, Daniel Buschek, Lior Zalmanson, and Mor Naaman. 2023. Co-Writing with Opinionated Language Models Affects Users' Views. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. Association for Computing Machinery, New York, NY, USA, 1–15. doi:10.1145/3544548.3581196
- [21] Maurice Jakesch, Megan French, Xiao Ma, Jeffrey T. Hancock, and Mor Naaman. 2019. AI-Mediated Communication: How the Perception That Profile Text Was Written by AI Affects Trustworthiness. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems (CHI '19)*. Association for Computing Machinery, New York, NY, USA, 1–13. doi:10.1145/3290605.3300469
- [22] Irving L. (Irving Lester) Janis. 1982. *Groupthink: Psychological Studies of Policy Decisions and Fiascoes*. Boston: Houghton Mifflin.
- [23] Tatsuya Kameda and Shinkichi Sugimori. 1993. Psychological Entrapment in Group Decision Making: An Assigned Decision Rule and a Groupthink Phenomenon. *Journal of Personality and Social Psychology* 65, 2 (1993), 282–292. doi:10.1037/0022-3514.65.2.282
- [24] Herbert C. Kelman. 1958. Compliance, Identification, and Internalization Three Processes of Attitude Change. *Journal of Conflict Resolution* 2, 1 (March 1958), 51–60. doi:10.1177/002200275800200106
- [25] Michael Knoll, Martin Götz, Elisa Adriasola, Amer Ali Al-Atwi, Alicia Arenas, Kokou A. Atitsogbe, Stephen Barrett, Anindo Bhattacharjee, Norman D. Blanco, Sabina Bogilović, Grégoire Bollmann, Janine Bosak, Cagri Bulut, Madeline Carter, Matej Černe, Susanna L. M. Chui, Donatella Di Marco, Gesa S. Duden, Vicki Elsey, Makoto Fujimura, Paola Gatti, Chiara Ghislieri, Steffen R. Giessner, Kenta Hino, Joeri Hofmans, Thomas S. Jönsson, Pazambadi Kazimna, Kevin B. Lowe, Juliana Malagón, Hassan Mohebbi, Anthony Montgomery, Lucas Monzani, Anne Nederveen Pieterse, Muhammed Ngoma, Emir Ozeren, Deirdre O'Shea, Christina Lundsgaard Ottsen, Jennifer Pickett, Anna A. Rangkuti, Sylwiusz Retowski, Farzad Sattari Ardabili, Razia Shaukat, Sílvia A. Silva, Ana Šimunić, Niklas K. Steffens, Faniya Sultanova, Daria Szűcs, Susana M. Tavares, Arun Tipandjan, Rolf van Dick, Dimitri Vasiljevic, Sut I. Wong, and Hannes Zacher. 2021. International Differences in Employee Motives: Scale Validation, Prevalence, and Relationships with Culture Characteristics across 33 Countries. *Journal of Organizational Behavior* 42, 5 (2021), 619–648. doi:10.1002/job.2512
- [26] Soohwan Lee, Seoyeong Hwang, Dajung Kim, and Kyungho Lee. 2025. Conversational Agents as Catalysts for Critical Thinking: Challenging Social Influence in Group Decision-making. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. Association for Computing Machinery, New York, NY, USA, 1–12. doi:10.1145/3706599.3719792
- [27] Soohwan Lee, Seoyeong Hwang, Mingyu Kim, Dajung Kim, and Kyungho Lee. 2026. Investigating LLM-Powered Dissenting Minority Support in Power-Imbalanced Group Decision-Making: Counterargument and Mediation as Intervention Strategies. *Proceedings of the ACM on Human-Computer Interaction* 10, 6 (Sept. 2026), CSCW106:1–CSCW106:40. doi:10.1145/3816954
- [28] Soohwan Lee, Seoyeong Hwang, and Kyungho Lee. 2025. Beyond Individual UX: Defining Group Experience(GX) as a New Paradigm for Group-centered AI. In *Companion Publication of the 2025 ACM Designing Interactive Systems Conference*. Association for Computing Machinery, New York, NY, USA, 357–362. doi:10.1145/3715668.3736348
- [29] SooHwan Lee, Mingyu Kim, Seoyeong Hwang, Dajung Kim, and Kyungho Lee. 2025. Amplifying Minority Voices: AI-Mediated Devil's Advocate System for Inclusive Group Decision-Making. In *Companion Proceedings of the 30th International Conference on Intelligent User Interfaces (IUI '25 Companion)*. Association for Computing Machinery, New York, NY, USA, 17–21. doi:10.1145/3708557.3716334
- [30] Colin MacDougall and Frances Baum. 1997. The Devil's Advocate: A Strategy to Avoid Groupthink and Stimulate Discussion in Focus Groups. *Qualitative Health Research* 7, 4 (Nov. 1997), 532–541. doi:10.1177/104973239700700407
- [31] Hannah Mieczkowski, Jeffrey T. Hancock, Mor Naaman, Malte Jung, and Jess Hohenstein. 2021. AI-Mediated Communication: Language Use and Interpersonal Effects in a Referential Communication Task. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW1 (April 2021), 17:1–17:14. doi:10.1145/3449091
- [32] Elizabeth Wolfe Morrison and Frances J. Milliken. 2000. Organizational Silence: A Barrier to Change and Development in a Pluralistic World. *The Academy of Management Review* 25, 4 (2000), 706–725. doi:10.2307/259200 jstor:259200
- [33] Serge Moscovici and Elisabeth Lage. 1976. Studies in Social Influence III: Majority versus Minority Influence in a Group. *European Journal of Social Psychology* 6, 2 (1976), 149–174. doi:10.1002/ejsp.2420060202
- [34] Hiroki Nakano, Jo Takezawa, Fabrice Matulic, Chi-Lan Yang, and Koji Yatani. 2026. Understanding Reader Perception Shifts upon Disclosure of AI Authorship. In *Proceedings of the 31st International Conference on Intelligent User Interfaces (IUI '26)*. Association for Computing Machinery, New York, NY, USA, 2131–2146. doi:10.1145/3742413.3789076
- [35] Charlan Nemeth, Keith Brown, and John Rogers. 2001. Devil's Advocate versus Authentic Dissent: Stimulating Quantity and Quality. *European Journal of Social Psychology* 31, 6 (2001), 707–720. doi:10.1002/ejsp.58
- [36] Ritika Poddar, Rashmi Sinha, Mor Naaman, and Maurice Jakesch. 2023. AI Writing Assistants Influence Topic Choice in Self-Presentation. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems (CHI EA '23)*. Association for Computing Machinery, New York, NY, USA, 1–6. doi:10.1145/3544549.3585893
- [37] Zoe A. Purcell, Maurice Jakesch, Mengchen Dong, Anne-Marie Nussberger, and Nils Köbis. 2025. Writing with AI Boosts Trust-Building Efficiency. *iScience* 28, 12 (Dec. 2025). doi:10.1016/j.isci.2025.114092
- [38] Peinuan Qin, Chi-Lan Yang, Nattapat Boonprakong, Jingzhu Chen, Yugin Tan, and Yi-Chieh Lee. 2026. AI Personalization Paradox: Reading Highlights for Personalized AI-Assisted Writing Increases Engagement but Undermines Autonomy and Ownership. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI '26)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/3772318.3791502
- [39] S. D. Reicher, R. Spears, and T. Postmes. 1995. A Social Identity Model of Deindividuation Phenomena. *European Review of Social Psychology* 6, 1 (Jan. 1995), 161–198. doi:10.1080/1479277943000049
- [40] Ronald E Robertson, Alexandra Olteanu, Fernando Diaz, Milad Shokouhi, and Peter Bailey. 2021. "I Can't Reply with That": Characterizing Problematic Email

- Reply Suggestions. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI '21)*. Association for Computing Machinery, New York, NY, USA, 1–18. doi:10.1145/3411764.3445557
- [41] Siavosh Sahebi, Paul Formosa, and Sarah Bankins. 2026. The AI Penalty and Disclosure Paradox: Trust, Authenticity and Knowledge Uptake in AI-mediated Communication. *Computers in Human Behavior: Artificial Humans* 8 (May 2026), 100304. doi:10.1016/j.chbah.2026.100304
- [42] Stefan Schulz-Hardt, Marc Jochims, and Dieter Frey. 2002. Productive Conflict in Group Decision Making: Genuine and Contrived Dissent as Strategies to Counteract Biased Information Seeking. *Organizational Behavior and Human Decision Processes* 88, 2 (July 2002), 563–586. doi:10.1016/S0749-5978(02)00001-8
- [43] Joongi Shin, Michael A. Hedderich, Andrés Lucero, and Antti Oulasvirta. 2022. Chatbots Facilitating Consensus-Building in Asynchronous Co-Design. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology (UIST '22)*. Association for Computing Machinery, New York, NY, USA, 1–13. doi:10.1145/3526113.3545671
- [44] Ben Shneiderman. 2020. Bridging the Gap Between Ethics and Practice: Guidelines for Reliable, Safe, and Trustworthy Human-centered AI Systems. *ACM Transactions on Interactive Intelligent Systems* 10, 4 (2020), 26:1–26:31. doi:10.1145/3419764
- [45] Dimitri Staufer, Frank Pallas, and Bettina Berendt. 2024. Silencing the Risk, Not the Whistle: A Semi-automated Text Sanitization Tool for Mitigating the Risk of Whistleblower Re-Identification. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*. Association for Computing Machinery, New York, NY, USA, 733–745. doi:10.1145/3630106.3658936
- [46] Michael Henry Tessler, Michiel A. Bakker, Daniel Jarrett, Hannah Sheahan, Martin J. Chadwick, Raphael Koster, Georgina Evans, Lucy Campbell-Gillingham, Tatum Collins, David C. Parkes, Matthew Botvinick, and Christopher Summerfield. 2024. AI Can Help Humans Find Common Ground in Democratic Deliberation. *Science* 386, 6719 (Oct. 2024), eadq2852. doi:10.1126/science.adq2852
- [47] Qiaosi Wang, Ida Camacho, Shan Jing, and Ashok K. Goel. 2022. Understanding the Design Space of AI-Mediated Social Interaction in Online Learning: Challenges and Opportunities. *Proc. ACM Hum.-Comput. Interact.* 6, CSCW1 (April 2022), 130:1–130:26. doi:10.1145/3512977
- [48] Benhua Xu and Yuhang Gao. 2025. How Does Uncertainty Avoidance Affect Employee Voice Behavior through the Dual-Path Ways of Work Engagement and Psychological Security? *BMC Psychology* 13, 1 (March 2025), 281. doi:10.1186/s40359-025-02598-z
- [49] Chi-Lan Yang, Jing Li, Xuhui Chang, Jingshu Li, Koji Yatani, and Yi-Chieh Lee. 2026. Can AI Be a Social Buffer? Investigating the Effect of AI-assisted Cognitive Reappraisal and Narrative Perspectives on Managing Difficult Workplace Conversations over Email. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI '26)*. Association for Computing Machinery, New York, NY, USA, 1–19. doi:10.1145/3772318.3790331