

Where Does the Disagreement Go? Group Reflexivity and AI Accounts of Team Decision-Making*

Soohwan Lee
Department of Design
UNIST
Republic of Korea
soohwanlee@unist.ac.kr

Kyungho Lee
Department of Design
UNIST
Republic of Korea
kyungho@unist.ac.kr




Figure 1: Three cases in which an AI-produced account can reshape a group's public position. In each panel, group interaction on the left passes through an agent to the account on the right. In A and B the account loses something: a concern that received no resolution does not reach a recap that records agreement (A), and an idea introduced on a shared canvas and restated in speech is attached only to the later speaker (B). C runs the other way. The agent makes mixed statements readable as three group positions, but a concern shared across the groups falls between the boundaries it draws. A and B suggest that an account should show more; C is a counter-case in which showing more adds structure of its own.

Abstract

AI agents mediate not only individual contributions but also the process through which many contributions become a team position. This matters when an AI-generated account of a meeting reaches another team, a manager, or an organization. Such an account may be accurate sentence by sentence while obscuring how the group reached its position, whose contributions shaped it, and which differences remain unresolved. We develop this position through two scenarios in which an account loses something: apparent consensus that leaves a concern unresolved, and a contribution that moves across speech and shared artifacts. We then present a counter-case in which an agent makes stakeholder positions visible, and that visibility may harden the divisions it was meant to clarify. Taken together, the cases suggest that the design problem is not how much an account should show, but whether a group can examine the interpretation it carries. We therefore ask how agents might scaffold members in tracing contributions and differences before an account travels, and we consider evaluating such support by whether a group, including members who brought different perspectives, can explain how it reached its decision, rather than by whether most members approve the account.

CCS Concepts

• **Human-centered computing** → **Computer supported cooperative work**; *Collaborative interaction*; *Natural language interfaces*; *HCI theory, concepts and models*.

Keywords

agent-mediated collaboration, group reflexivity, contribution attribution, accountability, collective account, meeting summarization

1 Introduction

AI agents are becoming part of how teams write, discuss, and decide. A member may use an agent to refine an idea before sharing it. An agent may also combine several members' inputs into meeting notes, a decision record, or a document for another team. This second role is common in larger organizations, where teams often encounter one another through reports, tickets, and recaps rather than through the original discussion. An agent-produced account can therefore do more than shorten a meeting. It can become the position through which a team appears to others.

Prior work has examined AI-mediated communication, meeting summaries, collective statements, and tools that prompt teams to reflect on their process [2, 3, 5, 22]. These approaches show how AI can support communication, common-ground finding, and meeting review. Less attention has been paid to what happens when an agent-produced account becomes the evidence of how a group

*Position paper accepted to the workshop "Nobody Did This": Contribution, Originality, and Accountability in Agent-Mediated Collaboration [8] at the 2026 ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW '26), Salt Lake City, UT, USA, October 2026.

reached its position. A concise account may omit an unresolved concern, attach an idea to the person who repeated it rather than the person who introduced it, or merge several positions into a statement that no member can fully explain. We use *group-level contribution dissolution* as a provisional term for this loss of distinct contributions and interactional context within an agent-produced representation of a group. This loss matters because disagreement and minority influence can sometimes prompt divergent thinking or reduce biased information seeking [14, 20]. It does not follow that every minority view is useful or should receive more weight. Making group differences more visible may also strengthen group boundaries or polarization [21].

We explore *group reflexivity* as one way to respond to this problem: members examine their goals, interaction, and decisions and consider whether to adjust them [19]. Rather than asking an agent to provide a finished reflection, we consider how it might indirectly scaffold a group in tracing contributions, revisiting different perspectives and values, and deciding what its account should carry forward. The group would still interpret that support in relation to its task, history, and power relations. For the CSCW '26 "Nobody Did This" Workshop, this position extends the workshop's account of contribution dissolution from the formation of individual contributions to the formation of a collective team account. We contribute two scenarios in which an account loses a consequential difference, a counter-case in which making differences visible may create a new problem (Figure 1), and questions for studying AI as a scaffold for group reflexivity. These cases do not support a general rule that systems should show more disagreement. They instead foreground the tension between preserving distinct contributions and avoiding representations that give weak, extreme, or temporary positions undue force.

2 Related Work

2.1 Agent-Mediated Contribution and Collective Accounts

AI-mediated communication includes cases in which an AI system modifies, augments, or generates a message to support a person's communication goals [5]. This mediation complicates attribution because people may develop an idea through interaction with an agent rather than bring a complete idea to it. The loss is broader than a missing label: documentation can record that AI was used without recovering who formed a judgment or who can answer for it. Related systems also synthesize several positions into a collective statement. Such systems can help groups find common ground [22], and meeting recap systems can support different post-meeting needs, since a single fixed-length summary serves quick overviews and detailed review unevenly [2]. However, a collective account does more than shorten text; it selects which contributions become visible to people who were not present. This creates two related questions: attribution asks who contributed to the account, and representation asks why the account can stand for the group. The second becomes important when a recap enters a ticket, decision log, or management report and gains authority through that workflow. A recipient may inspect the final wording or an AI-use disclosure but still be unable to see which disagreement was closed, deferred, or omitted.

2.2 Meeting Analytics, Reflection, and Group Reflexivity

Meeting support is moving from recording outcomes toward prompting reflection on the process. MeetingCoach presented post-meeting behavioral feedback [18], and Are We On Track? compared ambient visualization with active questions about meeting goals, identifying trade-offs among timing, interruption, and user control [3]. Structured collaborative reflection can also help people move from individual accounts toward shared interpretation, although evidence depends on the setting and form of scaffolding [4, 16]. What such systems can measure is constrained, however, by what collaborative work leaves behind. Speaking time describes access to the floor, not whether a short question changed the decision; a count of notes describes production, not whether the notes were taken up. Research on AI feedback for meeting inclusion also suggests that a behavioral measure cannot be separated from organizational norms, since a system may favor one communication style, expose an individual, or conflict with hierarchy [7]. We therefore treat speech, artifacts, and interaction as possible evidence rather than as a contribution score, and would keep captured traces, agent-inferred relations, and group-confirmed interpretations distinct rather than collapsing them into one display.

These studies also do not imply that an AI interpretation is itself group reflection. Individual feedback may help one member reconsider their behavior, and a team-facing display may create awareness of a shared pattern, but group reflexivity requires an additional social step: members compare accounts, question what the representation privileges, and decide whether to revise a shared account or future action.

A related line of work has agents intervene in the discussion itself rather than report on it. Systems that voice underrepresented positions or argue against an emerging consensus can raise psychological safety and satisfaction in power-imbalanced groups [9, 12]. How the support is delivered matters, however: relaying a minority member's view through an agent raised their participation while lowering their psychological safety, whereas AI-generated counterarguments did not show that cost [10]. These results concern the moment a position is expressed. Our question is what happens afterward, when the discussion is compressed into an account that travels. The two are not interchangeable: an intervention that helps a minority position enter the conversation does not ensure the position survives into the record, and a group that heard an objection may still circulate a recap that omits it. This distinction motivates our focus. We explore a division of labor in which an agent organizes evidence and prompts review, while members interpret that support and decide what may leave the meeting as a collective position.

3 AI-Facilitated Group Reflexivity

We use *AI-facilitated group reflexivity* to describe a group process rather than an automated analysis. Instead of writing a finished reflection, an agent might return selected traces and questions that help members examine how they worked and decide what to do next. These prompts could show how an idea moved across speech and shared artifacts, which perspectives an account merged, or where a disagreement remained unresolved. Members would

interpret these prompts in relation to their task and decide whether to revise the account, revisit a decision, or take no action. This is an indirect form of scaffolding: the agent structures an occasion for reflection but does not supply the group’s conclusion. Prior systems show that short, structured reflection can support teams in examining shared practices, while the effect depends on how reflection is organized [4, 16]. We call the occasion a *reflective checkpoint*. It might occur before notes are circulated or before one team hands work to another. We present it as a direction for design and study, not as a validated procedure.

A possible checkpoint could distinguish three layers that a summary often merges. The first is the *evidence*: an utterance, note, edit, or decision trace. The second is the agent’s *interpretation*: for example, a proposed link between a note and a later proposal, or an inference that an objection received no response. The third is the group’s *disposition*: members may confirm, correct, reject, defer, or leave that interpretation unresolved. An interface might place a source excerpt beside an inferred link and let members add a short reason for its disposition. It might also group contributions by issue rather than rank people by activity. This structure would not recover a complete history. It would instead help members distinguish what the system captured, what it inferred, and what the group reviewed. Whether members would use the distinction, or read any displayed inference as established, remains an empirical question.

This proposal does not treat visibility as a benefit by default. A minority view may be mistaken, weak, or extreme, and a group may have a reason to reject it. Making differences more visible may also harden temporary positions into group identities, as our counter-case illustrates. The intended role of scaffolding is therefore to help a group trace and interpret contributions, perspectives, values, and unresolved tension, not to decide which member contributed most or which view should prevail. Member review would not remove the agent’s influence: selection, framing, and timing still shape attention. A manager may dominate the review, or most members may approve an account because reopening it is costly. AI-facilitated group reflexivity may provide another occasion for the group to notice and contextualize this mediation, but it would not make the agent neutral, equalize power, or guarantee a better outcome.

4 Two Scenarios and a Counter-Case

We use two scenarios to examine what an agent-produced account may lose, followed by a counter-case about what visibility may cost (Figure 1). The first concerns unresolved disagreement and the second contribution across people and media. The counter-case then considers an AI visualization of stakeholder positions, which makes differences visible rather than losing them. These are constructed cases for examining design tensions, not empirical findings, and the risks we raise in each are conjectures we would want to test.

4.1 An Unresolved Concern That Does Not Reach the Account

A project team must select an implementation plan before the end of a meeting (Figure 1-A). Most members support one plan.

One member briefly raises an accessibility risk, but the meeting moves on before the group addresses it. The AI recap states that the team agreed on the plan, and another department treats the matter as settled. The account has turned a concern that received no resolution into apparent consensus. A reflective checkpoint might return the concern with its source excerpt and ask the group to mark it as incorporated, resolved, rejected with a reason, deferred, or unresolved. The account could still report the selected plan while carrying forward an unresolved risk.

This support would not imply that the concern should prevail. The team may have good reasons to proceed: the risk may be misunderstood, already handled elsewhere, or minor relative to the cost of addressing it. An explicit rejection with a reason is a legitimate outcome, and a system that treats every retained objection as a problem would push groups toward reopening settled matters. The purpose of a recorded status is to distinguish deliberative treatment from omission, not to give every minority view equal weight. A durable status may also expose the member who raised the concern, especially when managers can inspect the record. Research would therefore need to examine who can see the trace and whether members can record a disposition without turning disagreement into individual performance data.

4.2 Contributions That the Transcript Does Not Capture

A product team ideates with physical sticky notes on a shared canvas while talking (Figure 1-B). Members contribute in different ways. One writes many sticky notes but speaks little. Another posts few notes but talks often, asks questions, and keeps the session moving. A junior member writes a short note that reframes the problem; a senior member later explains and extends it in speech, and the final plan adopts the refined version. The AI recap links the idea only to the later statement, and describes the session mainly through what was said. Where supporting traces are available, a reflective checkpoint might show a possible link among the note, the speech, and the adoption, then let members correct it or supply missing context. The account could distinguish proposing, refining, and facilitating without selecting one owner or converting the acts into a score.

The difficulty here is not only that credit moved to the wrong person. It is that contributions arrive in different modalities, and a summary built mainly from transcript text can render them unevenly. An audio recording does not capture who writes or rearranges a physical note, and a photograph may preserve its content without revealing its author or how the group developed it. A digital shared canvas may retain authorship and edit history, but these traces are not necessarily available in an in-person session. The account may therefore preserve the spoken explanation while losing the context in which the idea took shape. Posting an idea, restating someone else’s idea so the group can act on it, and asking the question that changes the framing are different acts, and none of them reduces to volume in one channel. Activity measures do not settle this either. A member may speak little because they are listening, because their input is not needed, or because the group does not take up what they offer, and speaking time alone cannot distinguish these conditions. Contribution lineage may provide more context by showing

whether an idea moved into later work, but it requires capturing speech, notes, and edits, and the same trace that helps a member recover recognition could let a manager audit individual output. Attribution may also be unwanted when an idea is contested. A design might therefore let members decline attribution as well as claim it, and studies should examine who benefits when lineage becomes visible within a hierarchy.

4.3 A Counter-Case: The Costs of Making Differences Visible

One response to the omissions in A and B is to make contributions and differences in perspective visible so that members can examine them together and decide how to proceed. This counter-case considers a visualization of stakeholder positions: it may support group reflection, yet the way it organizes differences may introduce problems of its own.

Several departments negotiate a policy that affects cost, employee workload, and customer access (Figure 1-C). An agent represents statements from their meetings as points in a shared space and groups nearby points into visible positions. Such analysis can trace how suggestions are prioritized, discarded, or carried into recommendations [17]. It may reveal distinctions that a single summary would hide. The representation also adds its own structure: mixed and changing statements may appear as stable clusters associated with each department, while a concern shared across groups may be split by the cluster boundaries.

Our concern with this case is that a readable map may invite members to defend a group label rather than revise a position. We offer this as a conjecture, not a finding. It is informed by work in adjacent settings: making group identity salient in computer-mediated communication has been associated with stronger movement toward a group norm [21], and whether a stance is presented as a majority or as an isolated dissent has been associated with different patterns of stance change [13]. Neither studied a stakeholder map, so we read them as reasons the question is worth asking rather than as support for an answer. A possible alternative is to keep the unit of representation at the level of claims and time, showing overlap and change before assigning a group profile. This could preserve ambiguity but also make the display harder to read. Which representation helps a group reason without fixing its divisions is an open question, and one we would want to settle empirically before recommending either.

Together, the cases do not support a general rule that an account should show more. In A and B, the agent removes structure by dropping a concern or a contribution path. In C, it adds structure by drawing boundaries around positions. Both operations interpret the group. The common design question is whether members can examine that interpretation and decide how it should inform the account that travels.

5 Supporting Group Interpretation

Provenance records alone are unlikely to recover the social conditions that made a contribution accountable. Group reflection may therefore need to offer more than a log. Members may need to correct a proposed link, explain why an item was excluded, or retain a disagreement in the account. What counts as relevant evidence

would likely vary by task, since people judge credit through several factors rather than one quantity [6]. Letting a team select those criteria may itself create pressure to perform for them. A short review step before cross-team circulation might help members interpret the account in relation to their situation, but it would not indicate consent by default. Hierarchy or time pressure could turn a review button into routine approval. Timing matters for the same reason: active reflection can prompt reconsideration but also interrupt, while passive feedback is less disruptive and easier to ignore [3]. A checkpoint may therefore suit a high-consequence commitment or a handoff between teams better than every agenda item.

How such a step is evaluated matters as much as how it is built. Because making differences visible may also harden them, measuring only how much information the intervention exposes would be insufficient. Measures built around individual experience may also miss what happens at the group level, which has been raised as a reason to develop group-level constructs and metrics alongside individual ones [11]. Evaluation may need to separate whether most members liked or approved the intervention, whether participation changed, whether the group—including members who brought different perspectives and values—could explain its decision, and whether the agent steered the result. These outcomes need not move together. Two experiments on LLM-facilitated deliberation report versions of this gap: facilitation raised how much information groups shared without a significant effect on the decision outcome [1], and participants preferred facilitated discussion and cited inclusivity as the reason even though measures of participation equity did not improve, while the facilitators shifted some allocations [15]. A reflective checkpoint could plausibly show the same pattern, changing what a group discusses without changing what it decides. An endorsed account that the group cannot explain would therefore provide limited evidence of successful reflection. Any gains would also need weighing against surveillance, interpersonal exposure, and delay.

6 Conclusion

Agents intervene not only when individuals form and share contributions, but also when many contributions become one team position. Our cases suggest that an agent shapes that position whether it omits a difference or draws the lines between differences, and that both involve interpretations the group may never examine. Review may not be worthwhile for routine work, and when it is proportionate to the consequences remains open. We do not propose preserving every contribution or delaying every decision. We suggest studying how agents might scaffold teams in tracing contributions and differences, interpreting them in context, and explaining how an account came to speak for the group.

Acknowledgments

This research was partially supported by a grant from the Korea Institute for Advancement of Technology (KIAT) funded by the Government of Korea (MOTIE) (P0025495, Establishment of Infrastructure for Integrated Utilization of Design Industry Data) and the Industrial Technology Innovation Program (RS-2025-25457764, Smart CMF Design Recommendation and Real-Time Collaborative

3D Simulation System) funded by the Ministry of Trade, Industry & Energy (MOTIE, Korea).

References

- [1] Mohammed Alsobay, David M. Rothschild, Jake M. Hofman, and Daniel G. Goldstein. 2025. Bringing Everyone to the Table: An Experimental Study of LLM-Facilitated Group Decision Making. doi:10.48550/arXiv.2508.08242 arXiv:2508.08242 [cs]
- [2] Sumit Asthana, Sagi Hilleli, Pengcheng He, and Aaron Halfaker. 2025. Summaries, Highlights, and Action Items: Design, Implementation and Evaluation of an LLM-powered Meeting Recap System. *Proceedings of the ACM on Human-Computer Interaction* 9, 2 (May 2025), CSCW176:1–CSCW176:29. doi:10.1145/3711074
- [3] Xinyue Chen, Lev Tankelevitch, Rishi Vanukuru, Ava Elizabeth Scott, Payod Panda, and Sean Rintel. 2025. Are We On Track? AI-Assisted Active and Passive Goal Reflection During Meetings. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI ’25)*. Association for Computing Machinery, New York, NY, USA, 1–22. doi:10.1145/3706598.3714052
- [4] Jixiang Fan, Jiacheng Zhao, Sunggyeol Oh, Michael Bolmer, Yoonje Lee, Nick Flammer, Yuhao Chen, and D. Scott McCrickard. 2025. Structuring Collaborative Reflection: Integrating Diary Study and Focus Group Discussion. In *Companion Publication of the 2025 Conference on Computer-Supported Cooperative Work and Social Computing (CSCW Companion ’25)*. Association for Computing Machinery, New York, NY, USA, 244–248. doi:10.1145/3715070.3749233
- [5] Jeffrey T Hancock, Mor Naaman, and Karen Levy. 2020. AI-Mediated Communication: Definition, Research Agenda, and Ethical Considerations. *Journal of Computer-Mediated Communication* 25, 1 (March 2020), 89–100. doi:10.1093/jcmc/zmz022
- [6] Jessica He, Stephanie Houde, and Justin D. Weisz. 2025. Which Contributions Deserve Credit? Perceptions of Attribution in Human-AI Co-Creation. doi:10.48550/arXiv.2502.18357 arXiv:2502.18357 [cs]
- [7] Mo Houtti, Moyan Zhou, Daniel Runningen, Surabhi Sunil, Leor Porat, Harmanpreet Kaur, Loren Terveen, and Stevie Chancellor. 2026. Opportunities and Barriers for AI Feedback on Meeting Inclusion in Socioorganizational Teams. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI ’26)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/3772318.3791135
- [8] Kashif Imteyaz, Mohammad Rashidujaman Rifat, Divya Ramesh, Steven R. Rick, Simo Hosio, Hauke Sandhaus, Advait Sarkar, Christoph Riedl, and Saiph Savage. 2026. “Nobody Did This”: Contribution, Originality, and Accountability in Agent-Mediated Collaboration. doi:10.48550/arXiv.2607.26387 arXiv:2607.26387 [cs.CY] Workshop at the ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW ’26).
- [9] Soohwan Lee, Seoyeong Hwang, Dajung Kim, and Kyungho Lee. 2025. Conversational Agents as Catalysts for Critical Thinking: Challenging Social Influence in Group Decision-making. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA ’25)*. Association for Computing Machinery, New York, NY, USA, 1–12. doi:10.1145/3706599.3719792
- [10] Soohwan Lee, Seoyeong Hwang, Mingyu Kim, Dajung Kim, and Kyungho Lee. 2026. Investigating LLM-Powered Dissenting Minority Support in Power-Imbalanced Group Decision-Making: Counterargument and Mediation as Intervention Strategies. *Proceedings of the ACM on Human-Computer Interaction* 10, 6 (Sept. 2026), CSCW106:1–CSCW106:40. doi:10.1145/3816954
- [11] Soohwan Lee, Seoyeong Hwang, and Kyungho Lee. 2025. Beyond Individual UX: Defining Group Experience(GX) as a New Paradigm for Group-centered AI. In *Companion Publication of the 2025 ACM Designing Interactive Systems Conference (DIS ’25 Companion)*. Association for Computing Machinery, New York, NY, USA, 357–362. doi:10.1145/3715668.3736348
- [12] Soohwan Lee, Mingyu Kim, Seoyeong Hwang, Dajung Kim, and Kyungho Lee. 2025. Amplifying Minority Voices: AI-Mediated Devil’s Advocate System for Inclusive Group Decision-Making. In *Companion Proceedings of the 30th International Conference on Intelligent User Interfaces (IUI ’25 Companion)*. Association for Computing Machinery, New York, NY, USA, 17–21. doi:10.1145/3708557.3716334
- [13] Soohwan Lee and Kyungho Lee. 2026. Understanding Compliance and Conversion Dynamics in Multi-Agent Collectives. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI ’26)*. Association for Computing Machinery, New York, NY, USA, 1–26. doi:10.1145/3772318.3790385
- [14] Charlan Jeanne Nemeth and Joel Wachtler. 1983. Creative Problem Solving as a Result of Majority vs Minority Influence. *European Journal of Social Psychology* 13, 1 (1983), 45–55. doi:10.1002/ejsp.2420130103
- [15] Aaron T Parisi, Nithum Thain, Alden Hallak, Vivian Tsai, and Crystal Qian. 2026. Real-Time Group Dynamics with LLM Facilitation: Evidence from a Charity Allocation Task. In *Proceedings of the 2026 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’26)*. Association for Computing Machinery, New York, NY, USA, 25–59. doi:10.1145/3805689.3812253
- [16] Soya Park and Chinmay Kulkarni. 2023. Retrospector: Rapid Collaborative Reflection to Improve Collaborative Practices. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW2 (Oct. 2023), 293:1–293:20. doi:10.1145/3610084
- [17] Elinor Poole-Dayana, Deb Roy, and Jad Kabbara. 2025. An AI-Powered Framework for Analyzing Collective Idea Evolution in Deliberative Assemblies. doi:10.48550/arXiv.2509.12577 arXiv:2509.12577 [cs.CY]
- [18] Samiha Samrose, Daniel McDuff, Robert Sim, Jina Suh, Kael Rowan, Javier Hernandez, Sean Rintel, Kevin Moynihan, and Mary Czerwinski. 2021. MeetingCoach: An Intelligent Dashboard for Supporting Effective & Inclusive Meetings. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI ’21)*. Association for Computing Machinery, New York, NY, USA, 1–13. doi:10.1145/3411764.3445615
- [19] Michaëla C. Schippers, Michael A. West, and Amy C. Edmondson. 2017. Team Reflexivity and Innovation. *The Wiley Blackwell Handbook of the Psychology of Team Working and Collaborative Processes* (March 2017), 459–478. doi:10.1002/9781118909997.ch20
- [20] Stefan Schulz-Hardt, Marc Jochims, and Dieter Frey. 2002. Productive Conflict in Group Decision Making: Genuine and Contrived Dissent as Strategies to Counteract Biased Information Seeking. *Organizational Behavior and Human Decision Processes* 88, 2 (July 2002), 563–586. doi:10.1016/S0749-5978(02)00001-8
- [21] Russell Spears, Martin Lea, and Stephen Lee. 1990. De-Individuation and Group Polarization in Computer-Mediated Communication. *British Journal of Social Psychology* 29, 2 (1990), 121–134. doi:10.1111/j.2044-8309.1990.tb00893.x
- [22] Michael Henry Tessler, Michiel A. Bakker, Daniel Jarrett, Hannah Sheahan, Martin J. Chadwick, Raphael Koster, Georgina Evans, Lucy Campbell-Gillingham, Tatum Collins, David C. Parkes, Matthew Botvinick, and Christopher Summerfield. 2024. AI Can Help Humans Find Common Ground in Democratic Deliberation. *Science* 386, 6719 (Oct. 2024), eadq2852. doi:10.1126/science.adq2852